

Liquid IT : Toward a better compromise
between development scalability and
performance scalability not definitive

Etienne Brodu

July 1, 2015

Abstract

TODO translate from below when ready

Résumé

Internet étend l'économie à une échelle spatiale et temporelle sans précédent. Il permet à chacun de mettre un service à disposition de milliards d'utilisateurs, en seulement quelques heures. La plupart des grands services actuels ont commencés comme de simples applications créées dans un garage par une poignée de personnes. C'est cette facilité à l'entrée qui a permis jusqu'à maintenant une telle croissance sur le web. Google, Facebook ou Twitter en sont les exemples les plus flagrants. Pendant le développement d'une application, il est important de suivre cette croissance, au risque de se faire rattraper par la concurrence. Le développement est guidé par les besoins en terme de fonctionnalités, afin de vérifier rapidement si le service peut satisfaire l'audience. Des langages tel que Ruby ou Java se sont imposés comme les langages du web, justement parce qu'ils permettent d'intégrer facilement de nouvelles fonctionnalités.

Si une application répond correctement aux besoins, elle atteindra de manière virale un nombre important d'utilisateurs. Son audience peut prendre plusieurs ordres de grandeurs en quelques jours seulement, ou même en quelques heures suivant comment elle est relayée. Une application est dite scalable si elle peut absorber ces augmentations d'audience. Or il est difficile pour une application dont le développement est guidé par les fonctionnalités d'être scalable.

Au moment où l'audience commence à devenir trop importante, il est nécessaire de modifier l'approche de développement de l'application. Le plus souvent cela implique de la réécrire complètement en utilisant des infrastructures scalables qui imposent des modèles de programmation et des API spécifiques. Cela représente une charge de travail conséquente et incertaine. De plus, l'équipe de développement doit concilier cette nouvelle approche de développement scalable, avec la demande en fonctionnalités. Aucun langage n'a clairement réussi le compromis entre ces deux objectifs.

Pour ces raisons, ce changement est un risque pour la pérennité de l'application. D'autant plus que le cadre économique accorde peu de marges d'erreurs, comme c'est le cas dans la plupart des start-up, mais également dans de plus grandes structures.

Mon travail consiste à tenter d'écarter ce risque dans une certaine mesure. Ma thèse se base sur les deux observations suivantes. D'une part, Javascript

est un langage qui a énormément gagné en popularité ces dernières années. Il est omniprésent sur les clients, et commence à s'imposer également sur les serveurs avec Node.js. Il a accumulé une communauté de développeur importante, et est l'environnement d'exécution le plus largement déployé. De ce fait, il se place comme le langage principal du web, détrônant Ruby ou Java. D'autre part, l'exécution de Javascript a la particularité de ressembler à un pipeline. La boucle événementielle de Javascript est un pipeline qui s'exécute sur un seul cœur pour profiter d'une mémoire globale. On observe le même flux de messages traversant cette boucle événementielle que dans un pipeline.

L'objectif de ma thèse est de permettre à des applications développées en Javascript d'être automatiquement transformées vers un pipeline d'exécutions repartis. Nous construisons un compilateur permettant d'identifier les fonctions de Javascript et de les isoler dans ce que nous appelons des Fluxions. Un conteneur qui peut exécuter une fonction à la réception d'un message, et envoyer des messages pour continuer le flux vers d'autres fluxions. Les fluxions étant indépendantes, elles peuvent être déplacées d'une machine à l'autre. En transformant automatiquement un programme Javascript en Fluxions, on le rend scalable, sans effort.

Contents

1	Introduction	4
2	Context and objectives	5
2.1	The Web as a platform	7
2.1.1	From operating systems to the web	7
2.1.2	The languages of the web	8
2.1.3	Explosion of Javascript popularity	9
2.1.3.1	In the beginning	9
2.1.3.2	Rising of the unpopular language	11
2.1.3.3	Current situation	13
2.2	Highly concurrent web servers	15
2.2.1	Concurrency	15
2.2.1.1	Scalability	15
2.2.1.2	Tow faces of concurrency	16
2.2.2	Technological shift	17
2.2.2.1	The case for global memory	18
2.2.2.2	Rupture	18
2.2.3	A problem of memory	19
2.2.3.1	Event-loop and multi-processes paradigms	19
2.2.3.2	Invariance	20
2.3	Equivalence	21
2.3.1	Real-time web services	21
2.3.2	Transformation	22
3	State of the art	23
3.1	Javascript	24
3.1.1	Overview of the language	24
3.1.1.1	Functions as First-Class citizens	24

3.1.1.2	Lexical Scoping	24
3.1.1.3	Closure	24
3.2	Concurrency	24
3.2.1	Two known concurrency model	24
3.2.1.1	Thread	24
3.2.1.2	Event	24
3.2.1.3	Orthogonal concepts	24
3.2.2	Differentiating characteristics	24
3.2.2.1	Scheduling	24
3.2.2.2	Coordination strategy	24
3.2.3	Turn-based programming	24
3.2.3.1	Event-loop	24
3.2.3.2	Promises	24
3.2.3.3	Generators	24
3.2.4	Message-passing / pipeline parallelism -> DataFlow programming ?	24
3.3	Scalability	24
3.3.1	Theories	24
3.3.1.1	Linear Scalability	24
3.3.1.2	Limited Scalability	24
3.3.1.3	Negative Scalability	24
3.3.2	Scalability outside computer science (only if I have time)	25
3.4	Framworks for web application distribution	25
3.4.1	Micro-batch processing	25
3.4.2	Stream Processing	25
3.5	Flow programming	25
3.5.1	Functional reactive programming	25
3.5.2	Flow-Based programming	25
3.6	Parallelizing compilers	25
3.7	Synthesis	25
4	Fluxion	26
4.1	Fluxionnal Compiler	26
4.1.1	Identification	26
4.1.1.1	Continuation and listeners	26
4.1.1.2	Dues	26
4.1.2	Isolation	26
4.1.2.1	Scope identification	26

4.1.2.2	Execution and variable propagation	26
4.1.3	distribution	26
4.2	Fluxionnal execution model	26
4.2.1	Fluxion encapsulation	27
4.2.1.1	Execution	27
4.2.1.2	Name	27
4.2.1.3	Memory	27
4.2.2	Messaging system	27
5	Evaluation	28
5.1	Due compiler	28
5.2	Fluxionnal compiler	28
5.3	Fluxionnal execution model	28
6	Conclusion	29
A	Language popularity	30
A.1	PopularitY of Programming Languages (PYPL)	30
A.2	TIOBE	31
A.3	Programming Language Popularity Chart	32
A.4	Black Duck Knowledge	32
A.5	Github	34
A.6	HackerNews Poll	34

Chapter 1

Introduction

TODO 5p

Chapter 2

Context and objectives

Contents

2.1	The Web as a platform	7
2.1.1	From operating systems to the web	7
2.1.2	The languages of the web	8
2.1.3	Explosion of Javascript popularity	9
2.1.3.1	In the beginning	9
2.1.3.2	Rising of the unpopular language	11
2.1.3.3	Current situation	13
2.2	Highly concurrent web servers	15
2.2.1	Concurrency	15
2.2.1.1	Scalability	15
2.2.1.2	Tow faces of concurrency	16
2.2.2	Technological shift	17
2.2.2.1	The case for global memory	18
2.2.2.2	Rupture	18
2.2.3	A problem of memory	19
2.2.3.1	Event-loop and multi-processes paradigms	19
2.2.3.2	Invariance	20
2.3	Equivalence	21
2.3.1	Real-time web services	21

2.3.2	Transformation	22
-------	--------------------------	----

2.1 The Web as a platform

2.1.1 From operating systems to the web

With the invention of electronic computing machine, appeared the market for software applications. This market is not limited by marginal production cost ; software being a virtual product, the production and distribution cost for another unit is virtually null. The market is limited by the platform a software can be deployed on. The bigger the platform, the wider the market. It is interesting economically both for the provider, and for the consumer to standardize and widen the platform. For the provider because it can increase its revenue by reaching more consumers. And for the consumers, because of the competition between providers feeding innovation.



However, the first platforms for software started as products. The first commodity computers and the operating systems are economic products. It made sense for their manufacturer to increase their market share, so as to increase their revenue. The economical incentive. Microsoft successfully took over the market of operating system in the 90s, and cross the border of the monopoly more than once. However, before the internet, the distribution of software copies is still limited by physical medium. It still takes time to burn a CD, or a floppy, and to bring it to the consumer's home.

Sir Tim Berners Lee invented the world wide web in 1989. It was initially intended to share scientific documents and results. But it eventually became the distribution medium of choice for virtual products, software included. And pushed the scalability of software distribution.

Similarly to operating systems, Web browsers are products. They exposed innovative features to try to increase their market share. Among these features is the ability to run scripts directly inside the browser. Browsers partially replaced operating system as the platform to run software. But these heterogeneity in features between web browsers worsen the user experience on the web. Eventually, web technologies became specification to be respected by all browsers, so as to assure an homogeneous experience across the web. The web became the platform. It allows to deploy and run software at unprecedented scales. With web services, or Software as a Service, the distribution medium of software is so transparent that owning a product to

have an easier access is no longer relevant.

The browser and the operating system are only the foundation. The specifications assure the unification of the web as a platform. We explore now the different language to write and deploy applications on the web. TODO This subsection needs reviews. The message is still not clear.

2.1.2 The languages of the web

TODO This subsection needs reviews, it is poorly written.

In the early 90's, at the same time the web started developing, most of the more popular programming languages were released. Python(1991), Ruby(1993), Java(1994), PHP(1995) and Javascript(1995). With Moore's law predicting exponential increase in hardware performance, the industry realized that development time is more expensive than hardware. Low-level language were replaced by higher-level language with worse performance, but easier to develop with, to decrease the development time.

Java, developed by Sun Microsystems, imposes itself early as a language of choice and never really decreased. The language is executed on a virtual machine, allowing to write an application once, and to deploy it on heterogeneous machines. Additionally, it presents many hardware abstractions, such as memory garbage collection, and references, instead of manual memory allocation and pointers. For these reasons, it is easier to develop with, but shows slower performance than lower level languages like C/C++. The software industry quickly adopted it as its main development language. It is currently the most cited language on StackOverflow, and the second used on Github. It is generally in the first place of all the language popularity indexes. However, the software industry wants stable and safe solutions. This prudence generally slows down Java evolution. TODO More arguments needed here. Java is overengineered, overtooled, too verbose, it is not dynamic enough : OOP is outdated, and it lacks lambdas, and so on ...

Python is the second best language for everything. It is a general purpose language, currently quite popular for data science. In 2003, the release of the Django web frameworks brought the language to the web development scene.

Ruby was confined in Japan and almost unknown to the world until the release of Rails in 2005. With the release of this web framework, Ruby took-off and is still in active use. TODO why Ruby could have replaced Java ? dynamic language, functional paradigm

PHP stands for Personal Home Page Tools. It was designed to design

web pages since its beginning. It is probably one of the easiest language to start web development. However, according to several language popularity indexes, it is on a slow decline since a few years. Because of its simplicity, it often fail to grow projects to industrial size, and is being replaced by languages that succeed at bringing a prototype to industrial standards, like Java, Ruby or Python.

Since a few years, Javascript is slowly becoming the main language for web development. It is everywhere. It is present in every browser, and on the server as well with Node.js. Because of this unavoidable position, it became fast (V8, ASM.js) and usable (ES6, ES7). Additionally, it is a target for LLVM, allowing many languages to compile to Javascript, strengthening again its omnipresent position. I argue in this thesis, that Javascript is the language of choice to bring a prototype to industrial standards. ~~TODO un peu rapide ici~~ ~~TODO I should note the distinction between Javascript client and server side.~~ There is no choice on the client, while there is a choice on the server. So the popularity on the client is not as significant than on the server. However, as I said, Javascript unavoidable position made its success and its performance.

2.1.3 Explosion of Javascript popularity

2.1.3.1 In the beginning

Javascript was created by Brendan Eich at Netscape around May 1995, and released to the public in September. The initial name of the project was Mocha, then LiveScript, the name Javascript was finally adopted to leverage the trend around Java. Indeed, Java was considered the hot new web programming language at this time. It was quickly adopted as default language for web servers development, and everybody was betting on pushing Java to the client as well. The history proved them wrong.

When Javascript was released in 1995, the world wide web was on the rise.¹ Browsers were emerging, and started a battle to show off the best features and user experience to attract the wider public.² Javascript was released to be one of these features. Microsoft released their browser Internet Explorer 3 in June 1996 with a concurrent implementation of Javascript. They changed the name to JScript, to avoid trademark conflict with Oracle

¹<http://www.internetlivestats.com/internet-users/>

²to get an idea of the web in 1997 : <http://1x-upon.com/>

Corporation, who owns the name Javascript. The differences between the two implementations made difficult for a script to be compatible to both. At the time, banners and signs started to appear on web pages to inform the user about the ideal web browser to use for the best experience. This competition was fragmenting the web.

Netscape submitted Javascript to Ecma International for standardization in November 1996 to stop this fragmentation. In June 1997, ECMA International released ECMA-262, the first specification of ECMAScript, the standard for Javascript. A standard to which all browser should refer for their implementations.

The initial release for this specification was designed in a rush. The version released in 1995 was finished within 10 days. Because of this precipitation, the language has often been considered poorly designed and unattractive. Moreover, Javascript was intended to be simple enough to attract unexperienced developers, by opposition to Java or C++, which targeted professional developers. For these reasons, Javascript started with a poor reputation among the developer community.

But things evolved drastically since. The success of Javascript is due to many factors ; maybe the most important of all is the View Source menu that reveals the complete source code of any web application. *The view source menu is the ultimate form of open source*³. It is the vector of the quick dissemination of source code to the community, which picks, emphasizes and reproduces the best techniques. This brought open source and collaborative development before github. ~~TODO neither open source nor collaborative development are the correct terms~~ Moreover, all modern web browsers now include a Javascript interpreter, making Javascript the most ubiquitous runtime in history [**Flanagan2006**].

When a language like Javascript is distributed freely with the tools to reproduce and experiment on every piece of code. When this distribution is carried during the expansion of the largest communication network in history. Then an entire generation seizes this opportunity to incrementally build and share the best tools they can. This collaboration is the reason for the popularity of Javascript on the Web.

³<http://blog.codinghorror.com/the-power-of-view-source/>

2.1.3.2 Rising of the unpopular language

TODO why is Javascript unpopular ? cite some blog post like : <https://wiki.theory.org/YourLanguage> and add many blog posts titles in a mozaïc

Javascript started as a programming language to implement short interactions on web pages. The best usage example was to validate some forms on the client before sending the request to the server. This situation hugely improved since the beginning of the language. Nowadays, there is a lot of web-based application replacing desktop applications, like mail client, word processor, music player, graphics editor.

There is now more software services released to the public as web-based application compared to desktop clients.

ECMA International released several version in the few years following the creation of Javascript. The first and second version, released in 1997 and 1998, brought minor revisions to the initial draft. However, the third version, released in the late 1999, contributed to give Javascript a more complete and solid foundation as a programming language. From this point on, the consideration for Javascript kept improving.

In 2005, James Jesse Garrett released *Ajax: A New Approach to Web Applications*, a white paper coining the term Ajax [Garrett2005]. This paper points the trend in using this technique, and explain the consequences on user experience. Ajax stands for Asynchronous Javascript And XML. It consists of using Javascript to dynamically request and refresh the content inside a web page. It has the advantage to avoid requesting a full page from the server. Javascript is not anymore confined to the realm of small user interactions on a terminal. It can be proactive and responsible for a bigger part in the whole system spanning from the server to the client. Indeed, this ability to react instantly to the user gave to developer the feature to develop richer applications inside the browser. At the time, the first web applications to use Ajax were Gmail, and Google maps⁴.

Around this time, the Javascript community started to emerge. The third version of ECMAScript had been released, and it was homogeneously supported in the browsers. However, the DOM, and the XMLHttpRequest method, two components on which AJAX relies, still present heterogeneous interfaces among browsers. Javascript framework were released with the goal to straighten the differences between browsers implementations. Proto-

⁴A more in-depth analysis of the history of Ajax, given by late Aaron Swartz <http://www.aaronsw.com/weblog/ajaxhistory>

type⁵ and DOJO⁶ are early famous examples, and later jQuery⁷ and underscore⁸. These frameworks are responsible in great part to the wide success of Javascript and of the web technologies.

In the meantime, in 2004, the Web Hypertext Application Technology Working Group⁹ was formed to work on the fifth version of the HTML standard. This new version provide new capabilities to web browsers, and a better integration with the native environment. It features geolocation, file API, web storage, canvas drawing element, audio and video capabilities, drag and drop, browser history manipulation, and many mores. It gave Javascript the missing interfaces to become a rich environment to develop rich application in the browser. The first public draft of HTML 5 was released in 2008, and the fifth version of ECMAScript was released in 2009. These two releases, ECMAScript 5 and HTML5, represent a mile-stone in the development of web-based applications. Javascript became the programming language of this rising application platform.

Javascript, and web technologies are also used outside the web. NW.js¹⁰ and electron¹¹ are two solutions to deploy application built with web technologies. They use Node.js and Chromium. The Atom text editor¹², Popcorn Time¹³ and Light Table¹⁴ are example of such applications. However, if web applications are common choice for web service client on the desktop, HTML5 is not yet widely accepted as ready to build complete application on mobile, where performance and design are crucial. Indeed web-technologies are often not as capable, and well integrated as native technologies. But even for native development, Javascript seems to be a language of choice. An example is the React Native Framework¹⁵ from Facebook, which allow to use Javascript to develop native mobile applications. They prone the philosophy *"learn once, write anywhere"*, in opposition to the usual slogan *"write once,*

⁵<http://prototypejs.org/>

⁶<https://dojotoolkit.org/>

⁷<https://jquery.com/>

⁸<http://underscorejs.org/>

⁹<https://whatwg.org/>

¹⁰<https://github.com/nwjs/nw.js>

¹¹<https://github.com/atom/electron>

¹²<https://atom.io/>

¹³<https://popcorn.time.io/>

¹⁴<http://lighttable.com/>

¹⁵<https://facebook.github.io/react-native/>

run everywhere".¹⁶

2.1.3.3 Current situation

“When JavaScript was first introduced, I dismissed it as being not worth my attention. Much later, I took another look at it and discovered that hidden in the browser was an excellent programming language.”

—Douglas Crockford

The rise of Javascript is obvious on the web and particularly the open source communities. It also seems to be rising in the software industry. However, it is harder to give an accurate picture of the situation. The software industry is not as clear and open as the web. Moreover, there is no right metrics to accurately and directly measure programming language popularity. In the following paragraphs, I report some of the best metrics and indexes available freely on the web to try to represent the situation, both in the open source community and in the more opaque software industry. More detailed informations are available section A.

Available resources The TIOBE Programming Community index is a monthly indicator of the popularity of programming languages. Javascript ranks 6th on this index, as of April 2015, and it was the most rising language in 2014. It uses the number of results on many search engines as a measure of the popularity of a programming language. The results contains learning and training resources, forums logs, books and many other traces of the activity of a the community around the language. However, the measure used by the TIOBE is controversial, and might not be representative. It is a lagging indicator, and the number of pages doesn't represent the number of readers.

Alternatively, the PYPL index is based on Google trends to measure the number of requests on a programming language. Javascript ranks 7th on this index, as of May 2015. This index seems to be more accurate, as it depicts the actual interest of the community for a language. However, it is not representative as it only takes Google search into account.

From these indexes, the major programming languages are Java, C/C++ and C#. The three languages are still the most widely taught, and used to

¹⁶Used firstly by Sun for Java, but then stolen by many others

write softwares. But Javascript is rising to become an important language as well.


Developers collaboration platforms An indicator of the popularity and usage of a language is the number of developers and projects using it.

Github is the most important collaborative development platform, with around 9 millions users. Javascript is the most used language on github since mid-2011, with more than 320 000 repositories. The second language is Java with more than 220 000 repositories.

TODO : graph of Github repositories by languages


StackOverflow, is the most important Q&A platform for developers. It is a good representation of the activity around a language. Javascript is the second language showing the most activity on StackOverflow, with more than 840 000 questions. The first one is Java with more than 850 000 questions.

Black Duck Software helps companies streamline, safeguard, and manage their use of open source. For its activity, it analyzes 1 million repositories over various forges, and collaborative platforms to produce an index of the usage of programming language in open source communities. Javascript ranks second. C is first, and C++ third. Along with Java, the four first languages represent about 80% of all programming language usage.



TODO redo this graph, it is ugly.

Jobs The software industry is rather closed sourced, and its activity is rather opaque. All these previous metrics are representing the visible activity about programming language, but are not representative of the software industry. The trends on job openings gives an hint of the direction the software industry is heading towards. *Indeed* provide some trends over its database of job propositions. Javascript developers ranked at the third position, right after SQL and Java developers. Then come C# and C developers. This position means that Javascript might finally be on the edge to become a major language in the software industry, and become as important as Java and C/C++.



TODO redo this graph, it is ugly.

All these metrics in this section represent different faces of the current situation of Javascript adoption. With the rise of web applications, we can safely say that Javascript is one of most important language of this decade, alongside with Java and C/C++. It is widely used in open source projects, and everywhere on the web, as well as in the software industry.

2.2 Highly concurrent web servers

2.2.1 Concurrency

2.2.1.1 Scalability

The internet allows interconnection at an unprecedented scale. There is currently more than 16 billions devices connected to the internet, and it is

growing exponentially¹⁷. This massively interconnected network gives the ability for a web applications to be reached at the largest scale. A large web application like google search receives about 40 000 requests per seconds. That is about 3.5 billions requests a day¹⁸. This traffic is huge, but it remains roughly stable because the position of Google is popular and regularly used by a huge community of users. The spike of activity are hidden in the mass of average use.

However, the traffic of an unknown web application is much more uncertain. If the web application fits the market need, it might become viral when it is efficiently relayed in the media. As a concrete example, when a web application appears in the evening news, it might expect a huge spike in traffic. But most of the time, the spikes are unpredictable.

An unknown web application needs resources to follow the growth of its audience. This growth in resources might be steady enough to be planned, or it might be unexpected and challenging to meet. Scalability is the ability for a web application to adapt its load to the demand in a reasonable time. More precisely, an application is scalable, if the growth of its audience is proportional to the increase of its load on the resources.

Scalability mainly holds on the ability of a web application to respond to many concurrent requests. For this thesis, I reduce scalability to concurrency. A server is highly scalable if it is highly concurrent. *TODO better explanation of this simplification* A highly concurrent server is able to manage a large number of simultaneous request. In the 2000s, the limit to reach was to process 10 thousands simultaneous connections with a single commodity machine¹⁹. Nowadays, in the 2010s, the limit is set at 10 millions simultaneous connections at roughly the same price²⁰. With the growing number of connected devices on the internet, concurrency is a very important property in the design of web servers.

2.2.1.2 Tow faces of concurrency

Concurrency is the ability for an application to make progress on several tasks at the same time. It can be achieved either with parallelism, or with time-slicing concurrency.

¹⁷<http://blogs.cisco.com/news/cisco-connections-counter>

¹⁸<http://www.internetlivestats.com/google-search-statistics/>

¹⁹<http://www.kegel.com/c10k.html>

²⁰<http://c10m.robertgraham.com/p/manifesto.html>

The term Parallelism refers to techniques to make programs faster by performing several computations in parallel. This requires hardware with multiple processing units. Such hardware is nowadays omnipresent.²¹ Parallelism improves performances for computation. However, it is necessary for the parallel tasks to rely on independent states. Otherwise, two tasks could modify a shared state simultaneously resulting in its corruption.

Concurrency can be achieved with a single processing unit as well. The executions of the different tasks are interleaved in time. As an example, it is used to give a multi-tasking ability to operating systems. It allows to stretch the computation on one, or many cores. It enables scalability. `TODO review that.`

In time-slicing concurrency, the concurrent tasks can benefit of a shared state with different compromises depending on the scheduling strategy. In this thesis I focus on the cooperative scheduling used in the event-loop. It allows for a truly global memory and seems to be one of the easiest way for developers to write concurrent programs efficiently. Indeed, I presented in the previous section the popularity of Javascript, which is often implemented on top of this scheduling strategy (DOM, Node.js).

The other main scheduling strategy is preemptive scheduling. It is used in most execution environment in conjunction with multi-threading. However, it is known to be hard to manage, and should be avoided except when true concurrency is needed in concert with true shared state. Shared state could probably always be emulated with isolated memory and message passing.

2.2.2 Technological shift

Around 2004, the so-called Power Wall was reached. The clock of CPU is stuck at 3GHz because of the inability to dissipate the heat generated at higher frequencies. Additionally, the instruction-level parallelism is limited. Because of these limitations, a processor is limited in the number of instruction per second it can execute. Therefore, a coarser level of parallelism, like the task-level parallelism previously presented is the only option to achieve high concurrency. This parallelism requires the isolation of the memory of each independent task. This isolation is in contradiction with the best practices of software development. It creates a rupture between performance and development accessibility.

²¹https://wiki.haskell.org/Parallelism_vs._Concurrency

2.2.2.1 The case for global memory

The best practices in software development advocate to design a software into isolated modules. By following the best practices, the code base is split into modules. This modularity allows to understand each module of the application by itself, without an understanding of the rest. The understanding of the whole application emerges from the interconnections between the different modules. Modularity advocates three principles : encapsulation, a module contains the data, as well as the functions to manipulate this data ; separation of concerns, each module should have a clear scope of action, and it should not overlap with other scopes ; and loose coupling, each module should require no, or as little as possible knowledge about the definition of other modules. The main goal followed by these principles, is to help the developer to develop and maintain a large code-base.

Modularity is intended to avoid a different problem than the isolation required by parallelism. The former intends to avoid unintelligible spaghetti code ; while the latter avoids conflicting memory accesses resulting in corrupted state. The two goals are overlapping in the design of the application. Therefore, every language needs to provide a compromise between these two goals, and specialized in specific type of applications. TODO here I use language in both cases, it would be better to use a more generic term to refer to language or infrastructure

2.2.2.2 Rupture

Between the early development, and the maturation of a web application, the development needs are radically different. In its early development, a web application needs to quickly iterate over feedback from its users. The first reason of startup failures is the lack of market need²². The development reactivity is crucial. Therefore, the development team opt for a popular, and accessible language. The development team quickly releases a Minimum Viable Product as to get these feedbacks. “*Release early, release often*”, and “*Fail fast*” are the punchlines of the web entrepreneurial community.

As the application matures and its audience grows, the focus shift from the development speed to the scalability of the application. The development team shift from a modular language, to a language providing parallelism.

²²<https://www.cbinsights.com/blog/startup-failure-post-mortem/>

From this shift derives two problems. The first problem is the risk the development team needs to take to be able to grow the application. This risk usually implies for the development team to rewrite the code base to adapt it to a completely different paradigm, with imposed interfaces. It is hard for the development team to find the time, hence the money, or the competences to deploy this new paradigm. Indeed, the number two and three reasons for startup failures are running out of cash, and missing the right competences. The second problem is that after this shift, the development pace is different. Parallel languages are incompatible with the commonly learned design principles. The development team cannot react as quickly to user feedbacks as with the first paradigm.

This technological rupture is the demonstration that there is economically a need for a language that exposes a more sustainable paradigm to follow the evolution of a web application. A paradigm that it is easy to develop with, like in the beginning of a web application development, and yet parallelizable, so as to be scalable when the application matures. Such a paradigm would allow to develop with a global memory, so as to follow the best practices of software development, and be able to isolate the state of each concurrent tasks to parallelize them. Review this paragraph, the tow last sentences are almost equivalent, might need to merge them.

2.2.3 A problem of memory

It is easier to understand the possible parallelism of a cooking recipe than a computer program. That is because the modifications to the state are trivial in the cooking recipe, hence the interdependencies between operations. It is easy to understand that preheating the oven is independent from whipping up egg whites. While the interdependencies are not immediately obvious in a computer program. TODO is this metaphor useful here ? if yes, continue to a transition

2.2.3.1 Event-loop and multi-processes paradigms

We draw the similarity and differences between the event-loop and multi-processes paradigms.

The event-loop paradigm split the execution of a program into sequentially executed callbacks. All callbacks are executed sequentially so that each

callback as an exclusive access to the global memory storing the state of the application.

An application following the multi-processes paradigm splits the execution of a program into parallel processes. Each process has an exclusive access to its own memory to store a fraction of the state of the application. It uses message-passing to coordinate the state of the application at the global level, and to stream data from one process to the other.

Both paradigms encapsulate the execution, in callbacks or processes. Those containers are assured to have an exclusive access to the memory. The difference between these paradigms lies in the choice of exclusivity for the access to the state, and in the coordination of the state between the concurrent executions.

I focus on the coordination of state between the concurrent executions in concurrent programming. The main difference for the developer is how each paradigm assures to the developer the invariance in the memory state.

TODO it is not universal, but multi-process paradigms are also oriented around event-loops. An Event-loop is a multi-process on one machine. A multi-process is multiple event-loop running different part of the same program.

2.2.3.2 Invariance

I call invariance the assurance given to the developer that the global state of the application is not corrupted by the coordination between the concurrent executions. By corruption, I mean a modification of the state which was not intended by the developer. TODO not very clear

TODO the two following paragraphs might be better placed in the previous subsection In applications composed of multiple parallel processes, the memory of each processes is exclusive. The developer is aware that only the process can modify its memory. The processes propagate a modification to the state of the application by sending messages. Each process treat these messages one after the other. There is no risk of corrupted state by simultaneous, conflicting accesses. The invariance is explicit because the memory is isolated inside each process.

In application using cooperative scheduling on an event-loop, like Javascript, the memory is global. All the tasks, called events, are executed sequentially, so there is no possible simultaneous memory access. The developer is aware of the points in the code where the scheduler switches from one concurrent

execution to the other, so it can manage its state in atomic modification.

TODO the following paragraph is out of scope, either introduce it to say so, or remove it Actually, threads are already a compromise in the invariance, but it is a bad one, as it imposes the developer to annotate the program for possible parallelism (locks and so on) In applications composed of multiple parallel threads, the memory is shared between all the threads. The developer is aware that because of the preemptive scheduling strategy, the shared memory states of the application can be modified at any time by any thread. To prevent conflicting accesses on the memory, the developer locks every shared memory state during a modification. The developer assures itself the invariance of the memory.

The invariance exposed by the isolated processes and the event-loop are similar. The developer defines sequence of instructions with atomic access to the memory. And in both paradigms, these sequences communicate by sending messages to each other. The difference lies in the isolation of this memory. TODO we already said the two were similar, get deeper into this similarity for the invariance TODO this paragraph needs review

TODO Insert this : the event-loop isolate the memory in different time slots, while parallelism isolate the memory in different space slots.

TODO lack some steps and a proper transition

2.3 Equivalence

I argue that the language should propose to the developer an abstraction to encourage the best practices of software development. Then a compiler, or the execution engine, can adapt this abstraction so as to leverage parallel architectures. So as to provide to the developer a usable, yet efficient compromise. We propose to find an equivalence between the invariance proposed by the event loop paradigm and the invariance proposed by the multi-processes paradigm.

2.3.1 Real-time web services

This equivalence intend not to be universal. It focuses on a precise class of applications : real-time Web applications processing stream of requests from users. Web applications are now written in a stream fashion. Indeed, as I

showed previously, the event loop is executed like a pipeline. Web services can be seen as pipelines processing streams of requests.

In a stream processing application, there is roughly two kinds of usage of the global memory : data and state. Naively, the data represent a communication channel between different point in the application space, and the state represents a communication channel between different instant in time. The data flow from stage to stage through the pipeline, and are never stored on any parts. The state, on the other hand, remains in the memory to impact the future behaviors of the application. State might be shared by several parts of the application.

There are different kinds of state dependencies in applications components leading to different kinds of parallelism. In this thesis I argue that it is possible to parallelize a real-time web applications written on an event-loop because the strong dependencies mainly remains within a closed number of concurrent executions.

2.3.2 Transformation

The equivalence between the invariance of the event-loop paradigm and the invariance of the multi-process paradigm involves two steps. The first step is to isolate the different regions

Second the coordination of these regions by message-passing.

This equivalence should allow a compiler to transform an event loop into several parallel processes communicating by messages.

With this compiler, it would be possible to express an application with a global, so as to follow the design principles of software development. And yet, the execution engine could adapt itself to any parallelism of the computing machine, from a single core, to a distributed cluster.

Chapter 3

State of the art

3.1 Javascript

3.1.1 Overview of the language

3.1.1.1 Functions as First-Class citizens

3.1.1.2 Lexical Scoping

3.1.1.3 Closure

3.2 Concurrency

3.2.1 Two known concurrency model

3.2.1.1 Thread

3.2.1.2 Event

3.2.1.3 Orthogonal concepts

3.2.2 Differentiating characteristics

3.2.2.1 Scheduling

3.2.2.2 Coordination strategy

3.2.3 Turn-based programming

3.2.3.1 Event-loop

3.2.3.2 Promises

3.2.3.3 Generators

24

3.2.4 Message-passing / pipeline parallelism -> DataFlow programming ?

3.3 Scalability

3.3.2 Scalability outside computer science (only if I have time)

If I have time, I would like to try to explain why scalability is at the core of material engagement and information theory, and is at the core of our universe : the propagation of Gravity wave is an example : it is impossible to scale

3.4 Frameworks for web application distribution

3.4.1 Micro-batch processing

3.4.2 Stream Processing

3.5 Flow programming

3.5.1 Functional reactive programming

3.5.2 Flow-Based programming

3.6 Parallelizing compilers

OpenMP and so on

3.7 Synthesis

There is no compiler focusing on event-loop based applications

Chapter 4

Fluxion

4.1 Fluxionnal Compiler

Some parts of this are already written in the first paper. It needs a lot additional explanations and rewritting

4.1.1 Identification

4.1.1.1 Continuation and listeners

4.1.1.2 Dues

4.1.2 Isolation

4.1.2.1 Scope identification

Scope leaking

4.1.2.2 Execution and variable propagation

4.1.3 distribution

4.2 Fluxionnal execution model

Everything here is already written in the first paper : flx-paper. It only needs to be rewritten

4.2.1 Fluxion encapsulation

4.2.1.1 Execution

4.2.1.2 Name

4.2.1.3 Memory

4.2.2 Messaging system

Chapter 5

Evaluation

5.1 Due compiler

5.2 Fluxionnal compiler

5.3 Fluxionnal execution model

Chapter 6

Conclusion

Appendix A

Language popularity

A.1 PopularitY of Programming Languages (PYPL)

¹ The PYPL index uses Google trends² as a leading indicator of the popularity of a programming language. It search for the trend for each programming language by counting the number of searches of this language and the word "tutorial".

PYPL for May 2015

¹<http://pypl.github.io/PYPL.html>

²<https://www.google.com/trends/>

Rank	Change	Language	Share	Trend
1		Java	24.1%	-0.9%
2		PHP	11.4%	-1.6%
3		Python	10.9%	+1.3%
4		C#	8.9%	-0.7%
5		C++	8.0%	-0.2%
6		C	7.6%	+0.2%
7		Javascript	7.1%	-0.6%
8		Objective-C	5.7%	-0.2%
9		Matlab	3.1%	+0.1%
10	2× ↑	R	2.8%	+0.7%
11	5× ↑	Swift	2.6%	+2.9%
12	1× ↓	Ruby	2.5%	+0.0%
13	3× ↓	Visual Basic	2.2%	-0.6%
14	1× ↓	VBA	1.5%	-0.1%
15	1× ↓	Perl	1.2%	-0.3%
16	1× ↓	lua	0.5%	-0.1%

A.2 TIOBE

3

The TIOBE index uses many search engines as an indicator of the current popularity of programming languages. It counts the number of pages each search engine finds when queried with the language name and the word "programming". This indicator indicates the number of resources available, and the discussions about a given programming language.

Javascript was the most rising language of 2014 in the TIOBE index.

TIOBE for April 2015

³<http://www.tiobe.com/index.php/content/paperinfo/tpci/index.html>

Apr 2015	Apr 2014	Change	Programming Language	Ratings	Change
1	2	↑	Java	16.041%	-1.31%
2	1	↓	C	15.745%	-1.89%
3	4	↑	C++	6.962%	+0.83%
4	3	↓	Objective-C	5.890%	-6.99%
5	5		C#	4.947%	+0.13%
6	9	↑	JavaScript	3.297%	+1.55%
7	7		PHP	3.009%	+0.24%
8	8		Python	2.690%	+0.70%
9	-	2× ↑	Visual Basic	2.199%	+2.20%

A.3 Programming Language Popularity Chart

4

The programming language popularity chart indicates the activity of a given language in the online communities. It uses two indicators to rank languages : the number of line changed in github of, and the number of questions tagged with a certain language.

Javascript is ranked number one in this index. The Javascript community is particularly active online, and in the open source.

indeed.com

A.4 Black Duck Knowledge

5

The black-duck, which analyze the usage of language on many forges, and collaborative hosts, rank Javascript number 2, after C, and with about the same usage as C++.

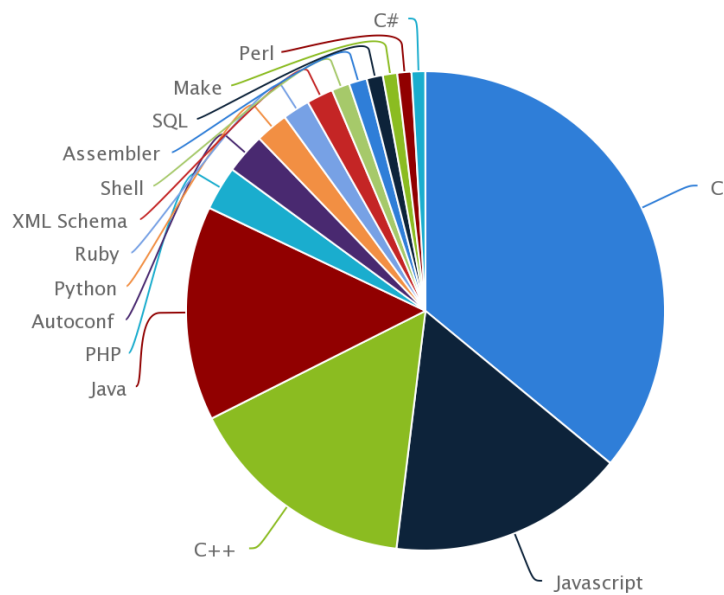
github.com sourceforge.net cpan.org rubyforge7.org planetsourcecode.com ddj.com

⁴<http://langpop.corger.nl>

⁵<https://www.blackducksoftware.com/resources/data/this-years-language-use>

Language	%
C	34.80
Javascript	15.45
C++	15.13
Java	14.02
PHP	2.87
Autoconf	2.65
Python	2.15
Ruby	1.77
XML Schema	1.73
Shell	1.18
Assembler	1.16
SQL	1.07
Make	0.94
Perl	0.92
C#	0.90

Releases within the last 12 months



Black Duck

A.5 Github

<http://github.info/>

A.6 HackerNews Poll

<https://news.ycombinator.com/item?id=3746692>

Language	Count
Python	3335
Ruby	1852
JavaScript	1530
C	1064
C#	907
PHP	719
Java	603
C++	587
Haskell	575
Clojure	480
CoffeeScript	381
Lisp	348
Objective C	341
Perl	341
Scala	255
Scheme	202
Other	195
Erlang	171
Lua	150
Smalltalk	130
Assembly	116
SQL	112
Actionscript	109
OCaml	88
Groovy	83
D	79
Shell	76
ColdFusion	51
Visual Basic	47
Delphi	45
Forth	41
Tcl	34
Ada	29
Pascal	28
Fortran	26
Rexx	13
Cobol	12

